

Assignment 1: House Regression

Introduction:

The dataset that was leveraged for this analysis was the House Prices Advanced Regression obtained from Kaggle. This dataset contains, “79 explanatory variables describing every aspect of residential homes in Ames, Iowa with the dependent variable being Sales Price to help predict the final price of each home. The goal is to establish and research the variables of house square footage and lot size that will impact the final price of the home.

Descriptive Statistics:

The dependent variable was identified as the Sales Price. By using this variable it was determined the average sale price of a home was 180,921.20. The minimum price was observed as 34,900 and the maximum was observed as 755,000. The standard deviation of this variable was shown to be 79,442.50. A histogram was constructed to display the distribution of the dependent variable. This histogram showed to be skewed heavily towards the right. While observing the histogram it displays that the bulk of the distribution sits within the lower price range(100,000-200,000).

Missing Data and Outliers:

The dataset includes housing features with numerous missing values, particularly in 'Alley', 'PoolQC', 'MiscFeature', and 'Fence', likely because not all properties have these attributes. While 'LotFrontage' and 'FireplaceQu' also show missing data, this may stem from recording issues. Such gaps in data can impact model accuracy and necessitate strategies like imputation or record removal. Also, a box plot analysis revealed outliers, with some properties' selling prices significantly deviating from the average, notably two examples exceeding \$700,000.

Total Lot Size and Total Square Foot:

The following predictors that were chosen to investigate were TotalBsmtSF, '1stFlrSF', '2ndFlrSF', 'LotFrontage', 'LotArea', 'WoodDeckSF', 'OpenPorchSF', and 'SalePrice'. The correlation was

calculated between these predictors and the sales price. It was determined that all calculated values show a positive correlation with the sales price variable. For example, a scatter plot was created between the sales price and the first floor sq ft which displays a positive correlation.

Merging predictor:

Two new variables were created to include the total square footage and total lot size. To create the total square footage variables such as 'TotalBsmtSF', '1stFlrSF', and '2ndFlrSF' were added to create 'TotalSF'. These variables, 'LotFrontage' and 'LotArea' were added up to make the total lot size variable. These combined variables potentially have an impact on the sales price. Multiple scatter plots were created to show the relationship between Sales Price, Total square footage, and Total Lot Size which displayed a largely positive correlation.

Min-max Scaling and Standard Scaling:

For our analysis, the 'SalePrice' variable underwent two normalization processes. The Min-Max scaling mapped the original prices into a new range where the smallest value corresponds to 0 and the largest to 1, transforming them into values like 0.2410. The standard scaling adjusted the prices so that their distribution has a mean of zero and a variance of one, resulting in normalized values such as -0.4885. These transformations are essential for models that are sensitive to the scale of data, ensuring that the variable's scale does not unduly influence the model's performance.

Conclusion:

Throughout this project, we gained valuable insights about the housing market, distribution of sales prices, missing data, outliers, and potential predictors. Additionally, using min-max scaling and standard scaling techniques can help identify if the variables can be used to create an advanced regression model. By using these findings it will provide a good source of knowledge to help predict the house prices as we seek to build a compelling model.

```
In [63]: import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
```

```
In [2]: house = pd.read_csv('train.csv')
house
```

```
Out[2]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl
...
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	Lvl
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl

1460 rows × 81 columns

Question 1

```
In [5]: house.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   Id                  1460 non-null   int64
 1   MSSubClass          1460 non-null   int64
 2   MSZoning            1460 non-null   object
 3   LotFrontage         1201 non-null   float64
 4   LotArea             1460 non-null   int64
 5   Street              1460 non-null   object
 6   Alley               91 non-null     object
 7   LotShape            1460 non-null   object
 8   LandContour         1460 non-null   object
 9   Utilities           1460 non-null   object
10   LotConfig           1460 non-null   object
11   LandSlope           1460 non-null   object
12   Neighborhood        1460 non-null   object
13   Condition1          1460 non-null   object
14   Condition2          1460 non-null   object
15   BldgType            1460 non-null   object
```

16	HouseStyle	1460	non-null	object
17	OverallQual	1460	non-null	int64
18	OverallCond	1460	non-null	int64
19	YearBuilt	1460	non-null	int64
20	YearRemodAdd	1460	non-null	int64
21	RoofStyle	1460	non-null	object
22	RoofMatl	1460	non-null	object
23	Exterior1st	1460	non-null	object
24	Exterior2nd	1460	non-null	object
25	MasVnrType	1452	non-null	object
26	MasVnrArea	1452	non-null	float64
27	ExterQual	1460	non-null	object
28	ExterCond	1460	non-null	object
29	Foundation	1460	non-null	object
30	BsmtQual	1423	non-null	object
31	BsmtCond	1423	non-null	object
32	BsmtExposure	1422	non-null	object
33	BsmtFinType1	1423	non-null	object
34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1422	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBlt	1379	non-null	float64
60	GarageFinish	1379	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1379	non-null	object
64	GarageCond	1379	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64
69	3SsnPorch	1460	non-null	int64
70	ScreenPorch	1460	non-null	int64
71	PoolArea	1460	non-null	int64
72	PoolQC	7	non-null	object
73	Fence	281	non-null	object
74	MiscFeature	54	non-null	object
75	MiscVal	1460	non-null	int64
76	MoSold	1460	non-null	int64
77	YrSold	1460	non-null	int64
78	SaleType	1460	non-null	object
79	SaleCondition	1460	non-null	object
80	SalePrice	1460	non-null	int64

```
dtypes: float64(3), int64(35), object(43)  
memory usage: 924.0+ KB
```

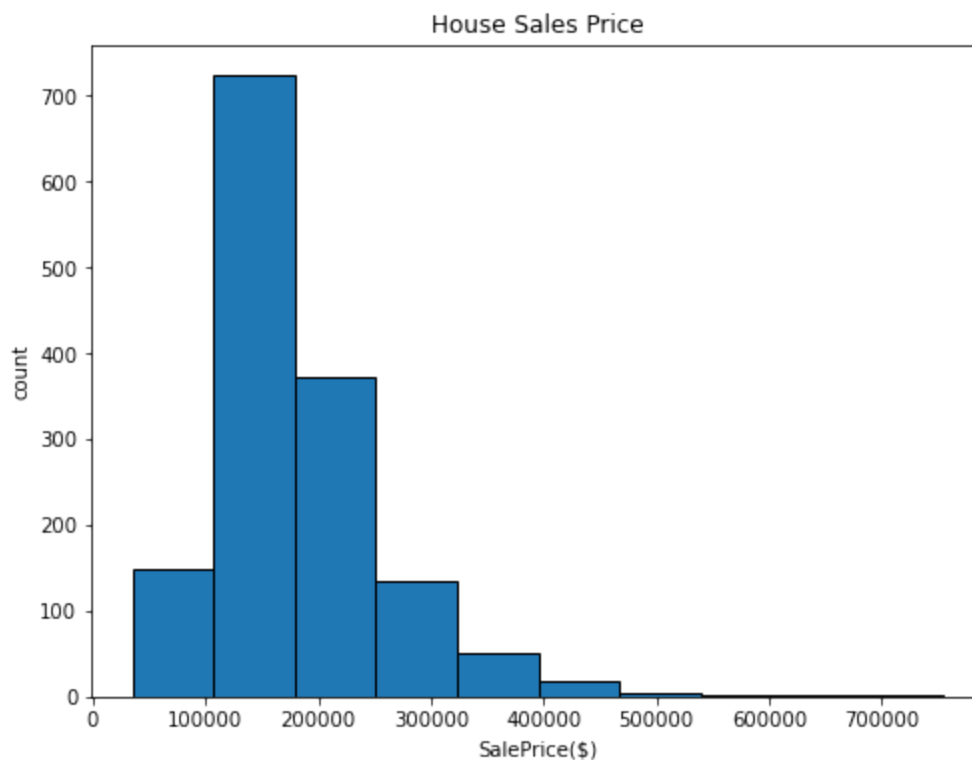
In [35]:

```
descriptive_stats = house['SalePrice'].describe()  
  
skewness = house['SalePrice'].skew()  
kurtosis = house['SalePrice'].kurtosis()  
  
print(descriptive_stats)  
print("Skewness:\n", skewness)  
print("Kurtosis:\n", kurtosis)
```

```
count      1460.000000  
mean       180921.195890  
std        79442.502883  
min        34900.000000  
25%       129975.000000  
50%       163000.000000  
75%       214000.000000  
max        755000.000000  
Name: SalePrice, dtype: float64  
Skewness:  
1.8828757597682129  
Kurtosis:  
6.536281860064529
```

In [65]:

```
plt.figure(figsize=(8,6))  
plt.hist(df['SalePrice'], edgecolor='black')  
plt.title('House Sales Price')  
plt.xlabel('SalePrice($)  
plt.ylabel('count')  
plt.show()
```



Question 2

In [64]:

```
house.isnull().sum()
missing_data = house.isnull().sum()
print("\nMissing data:")
print (missing_data[missing_data>0])
```

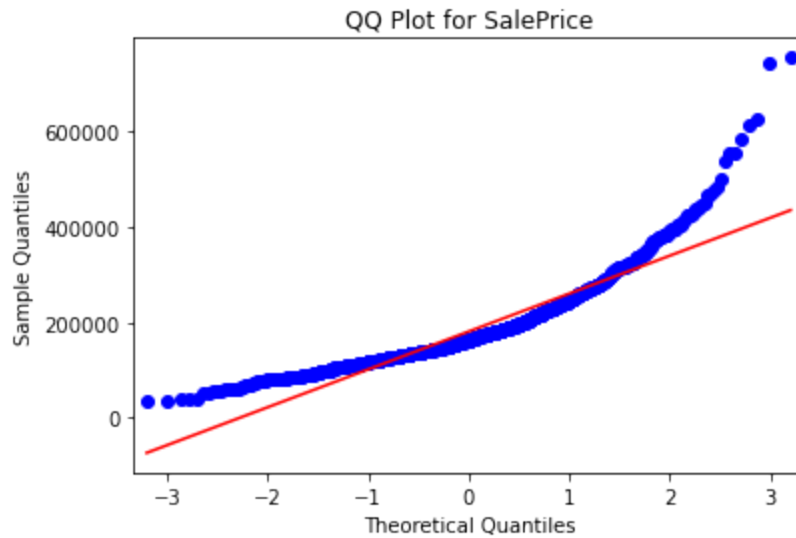
```
Missing data:
LotFrontage      259
Alley            1369
MasVnrType        8
MasVnrArea        8
BsmtQual         37
BsmtCond         37
BsmtExposure     38
BsmtFinType1     37
BsmtFinType2     38
Electrical        1
FireplaceQu      690
GarageType       81
GarageYrBlt      81
GarageFinish     81
GarageQual       81
GarageCond       81
PoolQC          1453
Fence            1179
MiscFeature      1406
TotalLotSize     259
dtype: int64
```

In [24]:


```
import matplotlib.pyplot as plt
import statsmodels.api as sm
from statsmodels.graphics.gofplots import qqplot

# Create the QQ plot for 'SalePrice'
plt.figure(figsize=(8, 6))
qqplot(house['SalePrice'], line='s') # 's' parameter adds a standardized line to
plt.title('QQ Plot for SalePrice')
plt.xlabel('Theoretical Quantiles')
plt.ylabel('Sample Quantiles')
plt.show()
```

<Figure size 576x432 with 0 Axes>



```
In [28]: plt.figure(figsize=(10 , 6))
sns.boxplot(x='SalePrice', data=df,
color='skyblue')
plt.show()
```



Question 3


```
In [85]: predict = house[['TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LotFrontage', 'LotArea', 'SalePrice']]
print(predict.corr())
```

	TotalBsmtSF	1stFlrSF	2ndFlrSF	LotFrontage	LotArea	SalePrice
TotalBsmtSF	1.000000	0.819530	-0.174512	0.392075	0.260833	0.613581
1stFlrSF	0.819530	1.000000	-0.202646	0.457181	0.299475	0.605852
2ndFlrSF	-0.174512	-0.202646	1.000000	0.080177	0.050986	0.319334
LotFrontage	0.392075	0.457181	0.080177	1.000000	0.426095	0.351799

LotArea	0.260833	0.299475	0.050986	0.426095	1.000000	0.263843
SalePrice	0.613581	0.605852	0.319334	0.351799	0.263843	1.000000

```
In [87]: predict.plot(x='SalePrice', y='1stFlrSF', style='o')
```

```
Out[87]: <AxesSubplot:xlabel='SalePrice'>
```



Question 4

```
In [77]: house['TotalSF'] = house['TotalBsmtSF'] + house['1stFlrSF'] + house['2ndFlrSF']
house.head()
```

```
Out[77]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilit
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllF
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllF
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllF
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllF
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllF

5 rows x 83 columns

```
In [78]: house['TotalLotSize'] = house['LotFrontage'] + house['LotArea']
house.head()
```

```
Out[78]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilit
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllF
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllF
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllF
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllF

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilit
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllF

5 rows × 83 columns

```
In [79]: df = house[['SalePrice', 'TotalSF', 'TotalLotSize']]
df
```


```
Out[79]:
```

	SalePrice	TotalSF	TotalLotSize
0	208500	2566	8515.0
1	181500	2524	9680.0
2	223500	2706	11318.0
3	140000	2473	9610.0
4	250000	3343	14344.0
...
1455	175000	2600	7979.0
1456	210000	3615	13260.0
1457	266500	3492	9108.0
1458	142125	2156	9785.0
1459	147500	2512	10012.0

1460 rows × 3 columns


```
In [81]: df.plot(x='SalePrice', y='TotalSF', style='o')
```

```
Out[81]: <AxesSubplot:xlabel='SalePrice'>
```



```
In [84]: df.plot(x='SalePrice', y='TotalLotSize', style='o')
```

Out [84]: <AxesSubplot:xlabel='SalePrice'>



In [52]: `print(df.corr())`

	SalePrice	TotalSF	TotalLotSize
SalePrice	1.000000	0.782260	0.319049
TotalSF	0.782260	1.000000	0.370958
TotalLotSize	0.319049	0.370958	1.000000

Question 5

In [62]:

```

from sklearn.preprocessing import MinMaxScaler, StandardScaler

sale_price = df['SalePrice'].values.reshape(-1, 1)

min_max_scaler = MinMaxScaler()
sale_price_minmax = min_max_scaler.fit_transform(sale_price)

standard_scaler = StandardScaler()
sale_price_standard = standard_scaler.fit_transform(sale_price)

print("Min-Max scaled SalePrice:\n", sale_price_minmax)
print("Standard scaled SalePrice:\n", sale_price_standard)

```

```

Min-Max scaled SalePrice:
[[0.24107763]
 [0.20358284]
 [0.26190807]
 ...
 [0.321622  ]
 [0.14890293]
 [0.15636717]]
Standard scaled SalePrice:
[[ 0.34727322]
 [ 0.00728832]
 [ 0.53615372]
 ...
 [ 1.07761115]
 [-0.48852299]
 [-0.42084081]]

```